

The Snake Eyes Paradox

Daniel M. Reeves
dreev@beeminder.com

2023-06-07 00:20

Problem Statement

You are offered a gamble. A pair of six-sided dice are rolled and unless they come up snake eyes you get a bajillion dollars. If they do come up snake eyes, you're devoured by snakes.

So far it sounds like you have a $1/36$ chance of dying, right?

Now the twist. First, I gather up an unlimited number of people willing to play the game, including you. I take 1 person from that pool and let them play. Then I take 2 people and have them play together, where they share a dice roll and either get the bajillion dollars each or both get devoured. Then I do the same with 4 people, and then 8, 16, and so on.

Eventually one of those groups will be devoured by snakes—hopefully not the group you're in—and then I stop. Is the probability that you'll die, given that you're chosen to play, still $1/36$?

Argument for NO: Due to the doubling, the final group of people that die is slightly bigger than all the surviving groups put together. So if you're chosen to play you have about a 50% chance of dying! 🤔🐍

Argument for YES: The dice rolls are independent and whenever you're chosen, what happened in earlier rounds is irrelevant. Your chances of death are the chances of snake eyes on your round: $1/36$. 😊

So which is it? What's your probability of dying, conditional on being chosen to play?

Some clarifications:

1. The game is not adversarial and the dice rolls are independent and truly random.
2. Choosing each group also happens uniformly randomly and without replacement.
3. The question is about the unrealistic case of an unbounded number of people but we can cap it and say that if no one has died after N rounds then the game ends and no one dies. We just need to then find the limit as N goes to infinity.
4. Importantly, in the finite version it's possible for no one to die. Just that the probability of that approaches zero as the size of the pool approaches infinity.
5. Again, we want the conditional probability that you die given that you're chosen to play. In other words, of the people chosen, what fraction die? In the unbounded case there's a 0% chance of infinitely many people being chosen and none dying; in the bounded case it's a tiny chance of a huge number chosen and none dying.

Don't turn the page until you're ready for the solution!

(Btw, this is only a veridical paradox. There's an unambiguously correct answer.)

Solution

We want the probability that you die given that you're chosen to play, $\Pr(\text{death} \mid \text{chosen})$. It seems like we can ignore the 0% chance of rolling not-snake-eyes forever and say that eventually about half the people who are chosen die, but let's Bayes it out carefully:

$$\begin{aligned}\Pr(\text{death} \mid \text{chosen}) &= \frac{\Pr(\text{chosen} \mid \text{death}) \Pr(\text{death})}{\Pr(\text{chosen})} \\ &= \frac{1 \cdot \Pr(\text{death})}{\Pr(\text{chosen})}.\end{aligned}$$

In the uncapped case, that conditional probability is undefined. You're part of an infinite pool so you have a 0% chance of being chosen and a 0% chance of dying. The probability we want is 0/0. **robot-with-smoke-coming-out-of-its-ears-emoji**

Since we can't directly calculate the probability in the infinite case, we have to take a limit.



To get a feel for where we're going, suppose you're 1 person in a huge but finite pool. Now suppose you're actually chosen. There are two ways that can happen:

1. The pool runs out and everyone survives.
2. The pool doesn't run out and you have about a 50% chance of dying.

But knowing that you're chosen is Bayesian evidence that we had many, many rounds of survival. If an early group died then most of the pool wasn't chosen, so probably you weren't chosen.

Thinking like a Bayesian means shifting your probability in light of evidence by seeing how surprised you'd be in various universes by that evidence. If an early group died then most people aren't chosen and in that universe you're surprised to be chosen. If *no* group died then everyone was chosen and in that universe you're fully unsurprised that you were chosen. That's the sense in which being chosen is Bayesian evidence that more people survived. In particular it's at least weak evidence that everyone survived.

So even with an absurdly huge pool of people, where there's *essentially* a 0% chance of everyone surviving, if you know you were chosen (which itself has near zero probability, but if) then that means you're more likely to be in that essentially-0%-probability universe where everyone survives.



Enough hand-waving and appeals to intuition. Let's Bayes it out to see what $\Pr(\text{death} \mid \text{chosen})$ is exactly, in the version where we stop after N rounds. Once we have that, we can take the limit as N goes to infinity.

First, let M be the size of the pool:

$$M = \sum_{i=0}^{N-1} 2^i = 2^N - 1.$$

And let p be the probability of snake eyes, $1/36$. We can now compute the probability of being chosen by summing up (1) the probability you're chosen for the first round, $1/M$, plus (2) the probability that the first group survives, $1 - p$, and that you're chosen for the 2nd round, $2/M$, plus (3) the probability that the first two groups survive and you're chosen for the 3rd round, etc. Writing that out as an equation gives this:

$$\begin{aligned}\Pr(\text{chosen}) &= \frac{1}{M} + (1 - p) \frac{2}{M} \\ &\quad + (1 - p)^2 \frac{4}{M} \\ &\quad + (1 - p)^3 \frac{8}{M} \\ &\quad + \dots \\ &\quad + (1 - p)^{N-1} \frac{2^{N-1}}{M} \\ &= \sum_{i=0}^{N-1} \frac{1}{M} 2^i (1 - p)^i.\end{aligned}$$

For $\Pr(\text{death})$ the calculation is very similar but every term is multiplied by p . To die, you have to be chosen and then roll snake eyes. This can happen on any round, all of which are mutually exclusive. We can then factor that p out and we have

$$\Pr(\text{death}) = p \cdot \Pr(\text{chosen}).$$

Working out that expression for $\Pr(\text{chosen})$ wasn't even necessary! We compute $\Pr(\text{death} \mid \text{chosen})$ like so:

$$\begin{aligned} \Pr(\text{death} \mid \text{chosen}) &= \frac{\Pr(\text{death})}{\Pr(\text{chosen})} \\ &= \frac{p \cdot \Pr(\text{chosen})}{\Pr(\text{chosen})} = p. \end{aligned}$$

It doesn't depend on N at all! The limit as N goes to infinity is just... p or $1/36$, the probability of rolling snake eyes. \square

Source: This is a variant of the Shooting-Room Paradox. Googling that yields a sea of confusion and wrongness and paywalled philosophy papers so I wrote this.

Acknowledgments

Thanks to Greta Goodwin, Bethany Soule, Christopher Moravec, and Stan Wagon, as well as Cesium-LifeJacket and Nix and others in the PEAR Discord for helpful discussion.

Thanks also to Manifold Markets where I [posed this question](#). The market converged to the correct answer amidst heated discussion in the comments:

Is the probability of dying in the Snake Eyes Paradox 1/36?

 Daniel Reeves 

 25  closes Jun 12 

98.1% chance

1D 1W 1M **ALL**

